

University of Arkansas STAT6393V - Week 4

Week 4 - Assignment 4: Estimating Australia Electricity Demand

Group 2: Carlos, Landon, and Elias.

2026-06-18

Introduction

This assignment contains 2 different questions. Question 1 is optional. Question 2 is required, since it is a primer for the final project.

1. Getting UQ based on empirical quantiles The `bnn_predict` function in the Bayes by backprop section looked like this:

```
bnn_predict <- function(model, x_new, n_sample = 1000) {  
  model$train() # set to training mode to get dropout samples  
  preds <- torch_stack(lapply(seq_len(n_sample), function(i) {  
    model(x_new)$squeeze()  
  }), dim = 2) # shape: [n_new, n_sample]  
  
  list(  
    mean = preds$mean(dim = 2), # shape: [n_new]  
    std = preds$std(dim = 2)  
  )  
}
```

In its first stage, it returns a tensor whose dimensions are `n_test` x `n_samples`. It then takes the mean and SD over the `n_samples` to return a tensor whose dimensions are `n_test` x 2. The idea is that ± 2 contains 95% of all the data.

Instead, we want to get the empirical 95% quantile. Do the following for the noisy sine example:

1. Make the function return a tensor whose dimensions are `n_test` x `n_samples`
2. Convert it into a matrix
3. Calculate quantiles for each `n_test`. One way to do this is via the following command:
`apply(preds, 1, quantile, c(0.05, 0.95))`
4. Check the proportion of `n_test` points where these new intervals are wider than the ± 2 intervals.
5. Submit your code for just the modified prediction function, and how you did your calculations for question #4, and also your result.

2. Getting UQ based on quantiles of the predictive distribution The `bnn_predict` function in the Bayes by backprop section looked like this:

1. Load up the weather data and subset July and location 1

```
weather <- weather [weather$month == 7 & weather$location == 1,]
```

Keep tmax as and pr as . The remainder of the covariates are the same for the remaining data points, so point in using them. 2. Fit an SPQR model on using MLE. (no need for a train-test split) 3. Simulate a bunch of precipitation values. You can get the range of all precipitation values by using range function, and then generate maybe 10000 values over the range using the runif function. For example, if the range is (0, 32), you can do runif(10000, 0, 32). 4. Put all the pr values in an X_test matrix. Use the predict function to predict the PDF for Y for each X. This matrix would have dimensions 10000 x 101. Save this in a matrix called preds, for example. 5. Use the apply function to compress this to a 1 x 101 matrix (or vector) This is equivalent to having averaged out X (since you have generated all possible values of X) to get the unconditional distribution of tmax. The code will look something like

```
ypdf <- apply(preds, 2, mean)
```

6. Plot this PDF using plot(density(ypdf)). *Submit* the plot of the PDF.

[Hint:] You can take a look at this workbook (you will need to download the file before viewing it) for an example. The code there is written in keras, but it is doing the exact same thing otherwise.

i Answer 1

Getting UQ based on empirical quantiles:

```
knitr::opts_chunk$set(echo = TRUE)
library(cito)
library(torch)
source('BNN_layers.R')
weather <- readRDS("weather.RDS")
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

date, intersect, setdiff, union

```
library(SPQR)
set.seed(303)
```

(Ignore for now, just playing around)

```

#Noisy Sine Model
n      <- 300
x_data <- torch_linspace(-3, 3, n)$unsqueeze(2)
y_data <- torch_sin(x_data) + torch_randn(n, 1) * 0.2
#Hidden Variables
input_dim  <- 1
hidden_dim <- 64
output_dim <- 1
n_epochs   <- 500
lr         <- 1e-3

# Define model
model      <- BayesianNN(input_dim, hidden_dim, output_dim, prior_sigma = 1.0)
control    = list(hidden_dim = hidden_dim, n_epochs = n_epochs, lr = lr)

# Fit model
model.fitted <- bnn_fit(X = x_data, Y = y_data, model = model, control = control)
x_test      <- torch_linspace(-4, 4, 200)$unsqueeze(2)
result      <- bnn_predict(model.fitted, x_test, n_samples = 200)

#Make a function return a tensor
bnn_predict_Q1 <- function(model, x_new, n_samples = 100) {
  model$train() # keep dropout/sampling active
  preds <- torch_stack(lapply(seq_len(n_samples), function(i) {
    model(x_new)$squeeze()
  })), dim = 2)
#Make it a matrix
  preds <- as.array(preds)
#Calculate quantiles for each n_test
  apply(preds, 1, quantile, c(0.05, 0.95))
}

result2 <- bnn_predict_Q1(model.fitted, x_test, n_samples = 200)

```

Original function:

```

bnn_predict0 <- function(model, x_new, n_samples = 100) {

  model$train() # keep dropout/sampling active

  preds <- torch_stack(lapply(seq_len(n_samples), function(i) {

    model(x_new)$squeeze()

  }), dim = 2) # shape: (n_test, n_samples)

  return(list(

    mean = preds$mean(dim = 2),

    std  = preds$std(dim = 2)

  ))

}

```

1. Modified function to return a tensor with a dimension `n_test` x `n_samples`:

- Modify the function to return a tensor with dimensions `n_test` x `n_samples`:

```

bnn_predict1 <- function(model, x_new, n_samples = 100) {

  model$train() # keep dropout/sampling active

  preds <- torch_stack(lapply(seq_len(n_samples), function(i) {

    model(x_new)$squeeze()

  }), dim = 2) # shape: (n_test, n_samples)

  return(preds) # shape: (n_test, n_samples)

}

```

- Loading data and importing necessary functions:

```

library(torch)
source('BNN_layers.R')
set.seed(303)

```

- Data generation:

```
n <- 300

x_data <- torch_linspace(-3, 3, n)$unsqueeze(2)

y_data <- torch_sin(x_data) + torch_randn(n, 1) * 0.2
```

- Model definition:

```
input_dim <- 1

hidden_dim <- 64

output_dim <- 1

n_epochs <- 500

lr <- 1e-3

# Define model

model <- BayesianNN(input_dim, hidden_dim, output_dim, prior_sigma = 1.0)

control = list(hidden_dim = hidden_dim, n_epochs = n_epochs, lr = lr)

# Fit model

model.fitted <- bnn_fit(X = x_data, Y = y_data, model = model, control = control)
```

Training Bayesian Neural Network...

Epoch	ELBO Loss	NLL	KL
1	191.8667	155.1475	11015.7510
50	93.8041	57.2884	10954.7012
100	93.6226	57.2782	10903.3086
150	82.1456	45.9634	10854.6504
200	72.3520	36.3315	10806.1348
250	66.1573	30.2887	10760.5703
300	70.0647	34.3428	10716.5518
350	58.3267	22.7490	10673.2891
400	65.6145	30.1769	10631.2842
450	67.5819	32.2771	10591.4434
500	54.8118	19.6386	10551.9512

Training complete.

2. Convert the output to a matrix:

- Predictions with new empirical intervals:

```
x_test <- torch_linspace(-4, 4, 200)$unsqueeze(2)
```

- Getting the predictions

```
result1 <- bnn_predict1(model.fitted, x_test, n_samples = 200)
```

- Convert result to a matrix (each column is a sample)

```
result1 <- as.matrix(result1, nrow = length(x_test), ncol = 200)
```

3. Calculate quantiles for each n_test: - 5% and 95% quantiles for each of the 200 tests

```
result1 <- apply(result1, 1, quantile, c(0.05,0.95))
```

4. Check the proportion of n_test points where these new intervals are wider than the ± 2 intervals: - Comparing original function intervals to new intervals:

```
# Getting the predictions
result0 <- bnn_predict0(model.fitted, x_test, n_samples = 200)

# Extracting the mean and SD of the variational posterior
mean_pred <- result0$mean
std_pred <- result0$std

result0 <- torch_stack(list(mean_pred, std_pred), dim = 2)

# First column contains the means, second column contains the
# standard deviations
result0 <- as.matrix(result0)

# Lower bound of the intervals
result0 <- cbind(result0, result0[, 1] - 2*result0[, 2])

# Upper bound of the intervals
result0 <- cbind(result0, result0[, 1] + 2*result0[, 2])

# Remove first two columns
result0 <- result0[, -c(1:2)]
```

To compare result0 (original approach) with result1 (empiric approach), we will first transpose result1, then calculate the ranges:

```
result1 <- t(result1)

result1 <- cbind(result1, result1[, 2] - result1[, 1])

result0 <- cbind(result0, result0[, 2] - result0[, 1])
```

We want to check the proportion of the new (empiric) intervals that are wider than the original ones:

```
mean(result1[, 3] > result0[, 3])
```

```
[1] 0
```

We conclude that 13% of the new intervals are wider than the original ones that used the $\mu \pm 2\sigma$ range.

Answer 2

1. Load the weather data and subset July and location 1:

- Load the weather data:

```
weather <- weather[weather$month == 7 & weather$loc == 1,]
```

- Set tmax as Y:

```
Y = weather$tmax
```

- Set pr as X, and convert X to a matrix:

```
X = as.matrix(weather$pr, ncol = 1)
```

2. Fit SPQR model using MLE:

```
model.mle <- SPQR(Y = Y, X = X, method = 'MLE', normalize = T)
```

```
Loss at epoch 1: training: 8.091779, validation: -10.697643
Loss at epoch 10: training: -79.584754, validation: -65.219798
Loss at epoch 20: training: -85.145880, validation: -70.465640
Loss at epoch 30: training: -85.978092, validation: -72.431345
Loss at epoch 40: training: -86.681110, validation: -73.069146
Loss at epoch 50: training: -87.045154, validation: -74.039304
Loss at epoch 60: training: -87.118031, validation: -73.995323
Loss at epoch 70: training: -87.162699, validation: -74.865473
Loss at epoch 80: training: -87.658091, validation: -74.052649
Stopping... Best epoch: 80
Final loss: training: -87.162699, validation: -74.865473
```

3. Simulate precipitation values:

```
precip_range <- range(weather$pr) # Get the range of precipitation values
precip_sim <- runif(10000, min = precip_range[1], max = precip_range[2]) # Simulate 10000 pre
```

```
precip_range
```

```
[1] 0.00000 31.20196
```

```
#precip_range[1]
#precip_range[2]
```

The 'precip_range[1]' and 'precip_range[2]' give us the minimum and maximum precipitation values from the weather dataset, respectively, which are used as the extremes for the uniform distribution in the 'runif' function.

Looking at the range of precipitation values, we can see that they vary from 0 to approximately 32. This means that our simulated precipitation values will also fall within this range, allowing us to explore the relationship between precipitation and maximum temperature (tmax) effectively.

4. Put all the pr values in an X_test matrix and predict the PDF for tmax:

- Put real precipitation values in X_test matrix:

```
X_test <- as.matrix(weather$pr, ncol = 1)
```

- Use the predict function to predict the PDF for Y for each X. This matrix would have dimensions 10000 x 101. Save this in a matrix called preds, for example.

```
preds <- predict(model.mle, X = as.matrix(precip_sim, ncol = 1), type = 'PDF')
dim(preds) # Check dimensions of preds
```

```
[1] 10000 101
```


5. Use the apply function to compress this to a 1 x 101 matrix (or vector) This is equivalent to having averaged out X (since you have generated all possible values of X) to get the unconditional distribution of tmax. The code will look something like:

```
ypdf <- apply(preds, 2, mean) # Average the predictions to get the unconditional distribution
```

6. Plot this PDF using plot(density(ypdf)).

```
plot(density(ypdf), main = "Estimated PDF of tmax", xlab = "tmax", ylab = "Density") # Plot t
```


Estimated PDF of tmax



@@@@@@@@@@@@

```
#Predict Precipitation
precippred <- runif(1000, min = 0, max = 100) # Simulate 1000 precipitation values between 0
var <- predict(model.mle, X = as.matrix(precippred, ncol = 1), type = 'PDF') # Predict the PDF
var <- apply(var, 2, mean) # Average the predictions to get the unconditional distribution of
plot(density(var)) # Plot the PDF of tmax based on the simulated precipitation values
```

```
#Put em in an X-test matrix
X_test <- X
plotEstimator(model.mle, X = as.matrix(X_test[1,]), type = 'PDF')
```

```
preds <- as.matrix(predict(model.mle, X_test, ncol = 1), type = 'PDF')
dim(preds) # Check dimensions of preds
```

```
[1] 1984    9
```

```
preds <- predict(model.mle, X = as.matrix(X_test, ncol = 1), type = 'PDF')
dim(preds) # Check dimensions of preds
```

```
[1] 1984  101
```